

A Fast Algorithm for Generating All Triangulations of Biconnected Plane Graphs

Tawkir Aziz Rahman

Department of Computer Science and Engineering
Bangladesh University of Engineering and Technology (BUET)
Dhaka-1000, Bangladesh

November 12, 2025

Abstract

In this paper, we address the problem of generating all triangulations of a biconnected plane graph. We present an algorithm for generating all triangulations of a biconnected plane graph G with n vertices. Our algorithm establishes a double recursive tree structure among all triangulations of G and generates each triangulation in $O(1)$ amortized time using $O(n)$ space. To the best of our knowledge, this is the first algorithm achieving $O(1)$ amortized time generation for biconnected plane graphs. The key innovations of our method are: (1) the concept of a *safe root vertex* for each face, ensuring the existence of a root triangulation that avoids multi-edges when combined with previously triangulated faces; (2) the use of local genealogical trees for each face, allowing independent processing of faces while maintaining global chord consistency; and (3) a dynamic global chord set that efficiently tracks and prevents multi-edge conflicts during the triangulation process. We provide proofs for the correctness, completeness, and efficiency of our algorithm, along with implementation details demonstrating its practical effectiveness.

1 Introduction

Triangulation of plane graphs is a fundamental problem in computational geometry with numerous applications, including VLSI floorplanning [5], general computational geometry [2], and graph drawing [7]. While computing a single triangulation of a plane graph can be performed efficiently, the exhaustive generation of *all* triangulations poses substantial computational challenges due to the combinatorial explosion of possible embeddings. In this paper, we focus on the problem of generating all triangulations of biconnected plane graphs—a problem that remains open despite significant advances on special subclasses of planar graphs.

1.1 Motivation and Challenges

The enumeration of all triangulations is critical for applications necessitating comprehensive exploration of the solution space, such as finding optimal triangulations under various cost metrics, studying structural properties of triangulation spaces, and analyzing the diameter of triangulation graphs. However, several fundamental difficulties make this problem challenging:

Exponential solution space. In the worst case, the number of triangulations for a given graph can be exponential in the number of vertices, rendering explicit storage and exhaustive

enumeration infeasible for large graphs. Practical algorithms must therefore generate triangulations on-the-fly without storing all of them simultaneously.

Duplicate avoidance. A major challenge in naive enumeration approaches is the frequent generation of identical triangulations multiple times. Traditional methods rely on storing previously generated triangulations and performing costly isomorphism checks, resulting in substantial memory overhead (potentially exponential) and at least logarithmic time complexity per triangulation for lookup operations.

Efficiency requirements. For enumeration algorithms, the benchmark for efficiency is output-sensitive time complexity, preferably constant or polynomial time per generated solution. This necessitates careful algorithm design to avoid redundant computations.

1.2 Related Work

Classical enumeration algorithms address the challenges of generating all triangulations through several strategies. The generate-and-filter method produces candidate objects and removes duplicates via filtering but is memory-intensive. Orderly generation [9] ensures canonical representations using sophisticated techniques, though it often involves expensive isomorphism testing. The reverse search paradigm [1] reduces memory requirements by implicitly defining a spanning tree over the solution space, enabling polynomial-space enumeration; however, its time complexity is typically $O(n^2)$ per triangulation in triangulation problems.

Early work by Hurtado and Noy [4] studied the triangulation graph for *convex polygons*, constructing a graph representing all triangulations and exploring its structural properties. Their approach requires generating all triangulations of smaller polygons as part of the process, potentially increasing complexity. For more general graphs, Li and Nakano [6] proposed an algorithm to enumerate all biconnected ‘based’ plane triangulations up to size n . Since their method incrementally builds from smaller graphs, it is inefficient when enumerating triangulations of a fixed-size graph.

The significant breakthrough for *triconnected* plane graphs came from Parvez, Rahman, and Nakano [8], who developed an elegant algorithm that achieves $O(1)$ time per triangulation using $O(n)$ space. Their key insight is to organize all triangulations into a tree structure, where parent-child relationships correspond to single edge flip operations. By carefully defining a unique *generating chord* for each triangulation, they ensure that flipping this chord produces a unique parent, thereby avoiding duplicates without explicit storage or isomorphism testing. This algorithm decomposes the graph into faces, triangulates each face independently, and efficiently combines the results.


However, this approach fundamentally relies on the graph being triconnected: in triconnected graphs, any two faces share edges only along their boundaries, preventing **multi-edge** conflicts. In general *biconnected* graphs, which may contain *separation pairs*, faces can share internal edges, leading to the multi-edge problem (see Section 1.3). This challenge obstructs the direct extension of the triconnected algorithm to all biconnected graphs.

Our work extends this framework to all biconnected graphs—a strictly more general class—while maintaining $O(1)$ amortized time and $O(n)$ space complexity. This is achieved by introducing novel techniques including safe root vertex selection, local genealogical trees, and dynamic global chord tracking to efficiently resolve multi-edge conflicts.

1.3 The Multi-Edge Problem in Biconnected Graphs

The Parvez-Rahman-Nakano algorithm critically relies on the graph being *triconnected*. Triconnectivity guarantees that any two faces share edges only along their boundaries—they cannot share internal chords. This property allows independent triangulation of faces without conflicts.

Triconnected Graph (Previous Work [8])



Biconnected Graph (This Work)




Figure 1: Comparison of graph classes. **Left:** Triconnected graphs (addressed by Parvez et al. [8]) have no separation pairs, so faces share only boundary edges. **Right:** Biconnected graphs (this work) may contain separation pairs (highlighted in red), allowing faces to share internal structure and creating multi-edge conflicts. Our algorithm extends the framework to this more general class through safe root selection and global chord tracking.

In contrast, general *biconnected* graphs, which have vertex connectivity exactly 2, do not satisfy this guarantee. Such graphs may contain *separation pairs*—pairs of vertices whose removal disconnects the graph. Faces separated by these pairs can share internal structure beyond their boundary edges. Consequently, independently triangulating faces can introduce **multi-edges**—multiple edges connecting the same pair of vertices—violating the simple graph property.

This **multi-edge problem** is the fundamental obstacle in extending the triconnected graph algorithm to biconnected graphs. Naively applying the face-by-face approach of Parvez et al. to biconnected graphs can produce invalid triangulations with parallel edges, making a direct extension non-trivial.


1.4 Our Contribution

In this paper, we present the first algorithm that generates all triangulations of biconnected plane graphs in $O(1)$ amortized time per triangulation using $O(n)$ space. Our algorithm successfully extends the tree-of-triangulations framework to the biconnected case by introducing novel techniques to resolve the multi-edge problem:

1. **Safe root vertex selection:** For each face, we introduce the concept of a *safe root vertex*—a vertex from which all initial chords can be drawn without creating multi-edges with previously triangulated faces. We prove that such vertices always exist and provide an efficient method to find them.
2. **Local genealogical trees:** Instead of a single global tree, we construct *local genealogical trees* for each face, organizing all valid triangulations of that face. These local trees are explored via depth-first search during recursive face processing.

Triconnected Graph

Face 1 triangulated



Biconnected Graph




Figure 2: The multi-edge problem in biconnected graphs. **Left:** In triconnected graphs, any two faces share only boundary edges, enabling independent triangulation. **Right:** In biconnected graphs with separation pairs, faces can share internal structure. Naive independent triangulation creates multi-edges (parallel edges between the same vertex pair), violating the simple graph property.

3. **Dynamic global chord tracking:** We maintain a *global chord set* that tracks all chords added across all faces during the recursive process. This enables efficient detection and prevention of multi-edge conflicts when triangulating each face.

We rigorously prove that our algorithm guarantees:

- **Completeness:** Every valid triangulation is generated exactly once (no duplicates, no omissions).
- **Correctness:** All generated triangulations are valid (no multi-edges, proper triangulation).
- **Efficiency:** $O(1)$ amortized time per triangulation with $O(n)$ space.

We also provide implementation details and experimental results demonstrating the algorithm's practical effectiveness on various graph instances.

2 Preliminaries

In this section, we establish the fundamental graph-theoretic concepts and notation used throughout this paper.

2.1 Basic Graph Theory

Let $G = (V, E)$ denote a connected simple graph with vertex set V and edge set E . We denote an edge connecting vertices v_i and v_j in V by (v_i, v_j) . The *degree* of a vertex v is the number of edges incident to v in G .

The *connectivity* $\kappa(G)$ of a graph G is the minimum number of vertices whose removal results in either a disconnected graph or a single-vertex graph. A graph is *k-connected* if $\kappa(G) \geq k$. In particular, a graph is *biconnected* if it is 2-connected (i.e., it has no cut vertices) and *triconnected* if it is 3-connected (i.e., it has no separation pairs).

2.2 Plane Graphs and Plane Embeddings

A graph G is *planar* if it can be embedded in the plane such that no two edges intersect except at their common endpoints. A *plane graph* is a planar graph with a fixed embedding in the plane. Such an embedding partitions the plane into connected regions called *faces*. The unbounded face is referred to as the *outer face*, while all other faces are *inner faces*.

A plane graph is called a *plane triangulation* if every face, including the outer face, is bounded by exactly three edges. In this paper, we allow edges to be represented by non-straight curves, consistent with standard practice in topological graph theory.

2.3 Cycles and Chords

A *cycle* C in a graph G is a sequence of distinct vertices $v_1, v_2, \dots, v_k$ such that the edges (v_i, v_{i+1}) exist for $1 \leq i < k$, and (v_k, v_1) closes the cycle. In a plane graph, a cycle divides the plane into two regions: the *inner region*, enclosed by C , and the *outer region*, which lies outside C . If the entire graph G is itself a cycle, both the inner and outer regions correspond to faces.

Consider a cycle C bounding a face F in a plane graph G , with vertices $v_1, v_2, \dots, v_n$ arranged in counterclockwise order. We denote this cycle as $C = \langle v_1, v_2, \dots, v_n \rangle$.

An edge (v_i, v_j) between two vertices on C is called a *chord* of C if it satisfies:

- $(v_i, v_j) \notin E(C)$, i.e., it is not an edge of the cycle itself, and
- the edge (v_i, v_j) lies entirely within the interior of the face F bounded by C .

Such a chord partitions the cycle C into two smaller cycles: one containing vertices $v_i, v_{i+1}, \dots, v_j$, and the other containing vertices $v_j, v_{j+1}, \dots, v_i$.

2.4 Triangulations

A *triangulation of a cycle* C is a maximal set of non-crossing chords that decomposes the inner region of C into triangles. In a triangulation T of C , the chord set is maximal in the sense that any chord not in T must intersect at least one chord in T . The sides of the resulting triangles are either chords from T or edges from the original cycle C .

For a cycle C with n vertices, any triangulation contains exactly $n - 3$ chords and produces exactly $n - 2$ triangular faces. We represent each triangulation T of C by the set of its chords. Given this chord set, the triangulation can be uniquely reconstructed.

A triangulation of a cycle C is called a *labeled triangulation* if the vertices of C are explicitly numbered sequentially from v_1 to v_n . If the vertices are not numbered, the triangulations are called *unlabeled triangulations*.

2.5 Visibility in Triangulations

In a triangulation T of a cycle C , we say that a vertex y is *visible from* a vertex x if there exists a chord (x, y) in T . This notion of visibility is fundamental to the structure of triangulations and plays a key role in our algorithm's construction of root triangulations.

2.6 Key Definitions for Our Algorithm

We now introduce several specialized definitions that are central to our algorithm for biconnected graphs.

Definition 1 (Multi-edge). A *multi-edge* occurs when two or more distinct edges connect the same pair of vertices in a graph. A *simple graph* contains no multi-edges or self-loops. All triangulations considered in this paper must be simple graphs.

Definition 2 (Separation Pair). A *separation pair* in a biconnected graph G is an unordered pair of vertices $\{u, v\}$ whose removal disconnects G into two or more components.

Definition 3 (Safe Root Vertex). A vertex r in a face F of a biconnected plane graph is called a *safe root vertex* if constructing a root triangulation from r does not create any multi-edge. Specifically, r is safe if for every non-neighbor vertex v of r on the face boundary, the chord (r, v) does not already exist in the global chord set S from previously triangulated faces.

Definition 4 (Genealogical Tree). A *genealogical tree* of triangulations for a cycle C is a rooted tree structure where:

- Each node represents a distinct triangulation of C ,
- The root is the root triangulation from a designated safe root vertex,
- Each edge represents a single chord flip operation that transforms one triangulation into another, and
- The parent-child relationship is defined by a unique generating chord.

Definition 5 (Global Chord Set). The *global chord set* is a dynamic data structure (typically implemented as a hash set) that maintains all chords currently present in the partially triangulated graph during the recursive enumeration process. This set enables constant-time detection and prevention of multi-edge formation.

Definition 6 (Hash Set). A *hash set* is a data structure implementing the set abstract data type using a hash table, providing average-case $O(1)$ time complexity for insertion, deletion, and membership testing operations.

Definition 7 (Leftmost Chord and Vertex Ordering). Let $C = \langle v_1, v_2, \dots, v_n \rangle$ be a cycle with vertices labeled in counterclockwise order. Chord (v_i, v_j) is **to the left of** chord (v_k, v_l) if: (1) $\min(i, j) < \min(k, l)$, OR (2) $\min(i, j) = \min(k, l)$ AND $\max(i, j) < \max(k, l)$. The **leftmost chord** is the chord to the left of all others.

Definition 8 (Edge Flip Notation). Given triangulation T with chord set C and potential chord $e = (v_a, v_b)$: Let e' be the unique chord in C intersecting e . We denote $\mathbf{T}(e) = T \setminus \{e'\} \cup \{e\}$ as the triangulation obtained by flipping. $\mathbf{T}_i(\mathbf{v}_1, \mathbf{v}_j)$ denotes the triangulation from T_i by flipping to create (v_1, v_j) .

These definitions provide the foundation for understanding our algorithm's approach to handling the multi-edge problem in biconnected graphs.

3 The Algorithm

3.1 Algorithm Overview

Our algorithm generates all triangulations of a biconnected plane graph G with k faces $F_1, F_2, \dots, F_k$ using a recursive depth-first search strategy. The approach consists of four key steps:

1. **Sequential face processing:** Process faces from F_1 to F_k in order
 2. **Local genealogical trees:** Construct a complete tree of valid triangulations for each face
 3. **Global conflict tracking:** Maintain a chord set to prevent multi-edge violations
 4. **Recursive composition:** Combine face triangulations to produce complete graph triangulations
-

3.2 Key Innovations

Our algorithm extends the Parvez-Rahman-Nakano framework from triconnected to biconnected graphs through three essential innovations:

3.2.1 Innovation 1: Local Genealogical Trees

Rather than constructing a single global genealogical tree (which fails for biconnected graphs due to multi-edge conflicts), we build **local genealogical trees** for each face independently.

Structure: Each face F_i has its own tree $\mathcal{T}(F_i)$ where:

- The root is a safe triangulation constructed from a safe root vertex
- Parent-child relationships are defined by edge flips within the face
- Trees are explored via DFS during recursive face processing


This decomposition enables independent triangulation of each face while maintaining global consistency through the chord set.

3.2.2 Innovation 2: Safe Root Vertex Selection


To ensure the root triangulation does not create multi-edges with previously triangulated faces, we introduce the concept of a **safe root vertex**.

Definition (Safe Root Vertex): A vertex r of face F_i is a *safe root vertex* if adding chords from r to all non-neighbor vertices of F_i does not conflict with any chord in the global set S (from previously triangulated faces $F_1, \dots, F_{i-1}$).

The safe root vertex guarantees a valid starting point for the local genealogical tree. We prove (Theorem X) that every face contains at least two safe root vertices, regardless of prior triangulations.



Process faces
 $F_1, F_2, \dots, F_k$
 sequentially




For each face,
 explore all
 local tree
 triangulations

Backtrack:
 Remove chords
 from S

Figure 4: DFS recursion structure with backtracking. The algorithm processes faces sequentially, exploring all triangulations in each face's local genealogical tree. The global chord set S is updated when entering a branch (add chords) and restored when backtracking (remove chords). Each path from root to leaf represents a unique complete triangulation of the graph.

Root from
 safe vertex r



Each edge flip
 creates child

Recursive
 descent

Local Genealogical Tree for Face F_i

Figure 5: Structure of a local genealogical tree for a single face. The root is constructed from a safe root vertex using a fan triangulation. Each child is obtained by flipping the generating chord (leftmost non-root chord). The tree structure ensures each triangulation appears exactly once, with unique parent-child relationships defined by edge flips.

3.2.3 Innovation 3: Dynamic Global Chord Management


The **global chord set** S tracks all chords in the current partial triangulation. It serves two critical purposes:

1. **Conflict detection:** Check if any chord already exists before adding it to a face triangulation
2. **Safe root identification:** Determine which vertices can serve as safe roots for the current face

The set is managed dynamically during DFS traversal:

- **Push:** When exploring a triangulation, add its chords to S
- **Recurse:** Process the next face with the updated S
- **Pop:** After exploring all descendants, remove chords from S (backtrack)

This ensures S always reflects exactly the chords in the current exploration path.



Dynamic management: Push chords when exploring, Pop when backtracking
 S always reflects current path in recursion tree

Figure 6: Dynamic global chord set management. As the algorithm explores the recursion tree via DFS, chords are added to S when entering a branch (push) and removed when backtracking (pop). This ensures S always contains exactly the chords from the current partial triangulation on the exploration path, enabling efficient multi-edge detection.


3.3 Core Concepts

3.3.1 Root Triangulation

Once a safe root vertex r is identified, we construct the **root triangulation** by adding chords from r to all non-neighbor vertices:

Definition (Root Triangulation): The *root triangulation* T_r of face F with safe root r is formed by adding all chords (r, v) for every vertex $v \in F$ that is not adjacent to r on the cycle boundary.

Root Triangulation from Safe Root $r = v_1$



Fan triangulation: Add all chords from r to non-neighbors

Figure 7: Construction of root triangulation from safe root vertex $r = v_1$ in a pentagon. The algorithm adds chords from r to all non-adjacent vertices (v_3 and v_4), creating a fan structure. Vertices v_2 and v_5 are already adjacent to r on the cycle boundary (blue edges), so no chords are needed. This construction guarantees no multi-edges if r is a safe root.

3.3.2 Local Genealogical Tree Structure

For each face, the local genealogical tree is defined by parent-child relationships based on edge flipping:

Definition 9 (Generating Chord). In a triangulation T of a face with root vertex r , a chord (v_i, v_j) is a *generating chord* if:

1. Neither v_i nor v_j is the root vertex r , and
2. It is the leftmost such chord (smallest index among non-root chords)

Definition 10 (Parent-Child Relationship). Let T be a triangulation with generating chord (v_i, v_j) forming triangles (v_i, v_j, v_p) and (v_i, v_j, v_q) on opposite sides. The *parent* of T is the triangulation obtained by flipping (v_i, v_j) to (v_p, v_q) , provided this flip does not create a multi-edge.

3.4 Algorithm Specification


Algorithm 1: Main Entry Point

Algorithm 1 StartTriangulation

- 1: **procedure** STARTTRIANGULATION(G, k)
 - 2: Label faces as $F_1, F_2, \dots, F_k$ arbitrarily
 - 3: Initialize global chord set $S \leftarrow \emptyset$
 - 4: $T \leftarrow \emptyset$ ▷ Empty partial triangulation
 - 5: FINDALLTRIANGULATIONS-cycle($F, T, 0$)
 - 6: **end procedure**
-

Current Triangulation T

Parent Triangulation



The parent of T is obtained by flipping the **generating chord**
(leftmost chord not incident to root v_1)

Figure 8: Parent-child relationship via generating chord. The generating chord is the leftmost chord in the triangulation that does not involve the root vertex. Flipping this chord produces the unique parent triangulation. This defines a tree structure where each non-root triangulation has exactly one parent.

Algorithm 2: Safe Root Construction

Algorithm 2 ConstructRootTriangulation

```

1: procedure CONSTRUCTROOTTRIANGULATION( $F$ , index)
2:    $r \leftarrow \text{FINDSAFEROOTVERTEX}(F, \text{index})$ 
3:    $T_r \leftarrow \emptyset$ 
4:   Let  $F = \langle v_1, v_2, \dots, v_n \rangle$  with  $v_r = \text{safe root vertex}$ 
5:
6:   for each vertex  $v_j$  in  $F$  do
7:     if  $v_j$  is not adjacent to  $v_r$  on cycle boundary then
8:       Add chord  $(v_r, v_j)$  to  $T_r$ 
9:     end if
10:  end for
11:
12:  return  $T_r$ 
13: end procedure

```

Algorithm 3 FindSafeRootVertex

```
1: procedure FINDSAFEROOTVERTEX( $F$ , index)
2:   Let  $F = \langle v_1, v_2, \dots, v_n \rangle$ 
3:   startIndex  $\leftarrow n - 1$ 
4:   endIndex  $\leftarrow 0$ 
5:
6:   while startIndex > endIndex + 1 do
7:     if chord  $(v_{\text{startIndex}}, v_{\text{endIndex}})$  exists in  $S$  then
8:       startIndex  $\leftarrow$  startIndex - 1
9:     else
10:      endIndex  $\leftarrow$  endIndex + 1
11:    end if
12:  end while
13:
14:  return  $v_{\text{endIndex}}$  ▷ Safe root vertex found
15: end procedure
```

Time Complexity: $\mathcal{O}(n)$ where n is the number of vertices in the face.

Algorithm 3: Face-Level Triangulation

Algorithm 4 FindAllTriangulationsCycle

```
1: procedure FINDALLTRIANGULATIONS CYCLE( $F, T, i$ )
2:    $n \leftarrow$  number of vertices of face  $F_i$ 
3:    $T_{ir} \leftarrow$  CONSTRUCTROOTTRIANGULATION( $F, i$ ) ▷  $\mathcal{O}(n)$ 
4:
5:   Add all chords of  $T_{ir}$  to  $S$  ▷  $\mathcal{O}(n)$ 
6:   OUTPUTTRIANGULATION( $T, T_{ir}, i$ )
7:    $T_i \leftarrow T_{ir}$ 
8:
9:   ▷ Start from  $n - 1$ : vertex  $v_n$  is adjacent to  $v_1$ , so  $(v_1, v_n)$  is an edge
10:  ▷ Similarly  $v_2$  is adjacent to  $v_1$ , so we start from  $j = 3$ 
11:  for  $j = n - 1$  down to 3 do
12:    Add chord  $(v_1, v_j)$  to  $S$ 
13:    FINDALLCHILDTRIANGULATIONS CYCLE( $T_i(v_1, v_j), T, i$ )
14:    Remove chord  $(v_1, v_j)$  from  $S$ 
15:  end for
16:
17:  Remove all chords of  $T_{ir}$  from  $S$  ▷  $\mathcal{O}(n)$ 
18: end procedure
```

Algorithm 4: Recursive Child Generation

Understanding the Critical Condition: The algorithm explores a child triangulation whenever *at least one* of these conditions holds:

- **Case 1:** (v_1, v_j) is already a chord in $T_i \rightarrow$ We can perform an edge flip without conflicts
- **Case 2:** $(v_1, v_j) \notin S \rightarrow$ Adding this chord will not create a multi-edge

- **Pruning case:** If neither holds (chord not in T_i **and** chord in S), we skip this branch to avoid multi-edges

Algorithm 5 FindAllChildTriangulationsCycle

```

1: procedure FINDALLCHILDTRIANGULATIONS-cycle( $T_i, T, i$ )
2:   OUTPUTTRIANGULATION( $T, T_i, i$ )
3:
4:   if  $T_i$  has no generating chord then
5:     return ▷ Leaf node reached
6:   end if
7:
8:   Let  $(v_b, v_{b'})$  be the leftmost generating chord of  $T_i$ 
9:
10:  for all  $j \geq b$  do
11:    ▷ Safe to explore: existing chord OR no conflict
12:    if  $(v_1, v_j)$  is a chord of  $T_i$  or  $(v_1, v_j) \notin S$  then
13:      Add chord  $(v_1, v_j)$  to  $S$ 
14:      FINDALLCHILDTRIANGULATIONS-cycle( $T_i(v_1, v_j), T, i$ )
15:      Remove chord  $(v_1, v_j)$  from  $S$ 
16:    end if
17:  end for
18: end procedure

```

Key Property: The algorithm explores *all* valid branches (every j satisfying the condition), ensuring completeness.

Algorithm 5: Output Control

Algorithm 6 OutputTriangulation

```

1: procedure OUTPUTTRIANGULATION( $T, T_i, i$ )
2:   if  $i = k$  then ▷ All faces triangulated
3:     output  $(T \cup T_i)$ 
4:   else
5:     FINDALLTRIANGULATIONS-cycle( $F, T \cup T_i, i + 1$ )
6:   end if
7: end procedure

```


3.5 The Edge Flip Mechanism

3.5.1 How Edge Flipping Works

An edge flip transforms one triangulation into another by replacing a chord:

1. **Select the chord:** Choose a chord (v_a, v_b) in triangulation T
2. **Identify adjacent triangles:** Find triangles (v_a, v_b, v_p) and (v_a, v_b, v_q) on opposite sides

3. **Perform the flip:** Remove (v_a, v_b) and add (v_p, v_q) to create T'
4. **Validate:** Verify $(v_p, v_q) \notin S$ to avoid multi-edges



Edge flip: Remove chord (v_a, v_b) and add chord (v_p, v_q)
Transforms one valid triangulation into another

Figure 9: Edge flip operation. A chord (v_a, v_b) forming two adjacent triangles is removed and replaced with the opposite diagonal (v_p, v_q) . This operation transforms one triangulation into another while preserving validity. Each child in the genealogical tree is obtained from its parent by a single edge flip.

Notation: We write $T(e)$ for the triangulation obtained by flipping to create chord e . For example, $T_i(v_1, v_j)$ denotes the triangulation created by flipping to form chord (v_1, v_j) .

3.5.2 Tree Structure Properties

The genealogical tree has these essential properties:

- **Unique parent:** Each triangulation (except root) has exactly one parent, determined by flipping its generating chord
- **Multiple children:** A triangulation may have several children, obtained by different valid flips
- **No cycles:** The tree structure (not a DAG) prevents any triangulation from appearing twice
- **Complete coverage:** Every valid triangulation is reachable from the root via edge flips

3.5.3 Completeness Through Tree Exploration

The genealogical tree guarantees we find all triangulations:

- **Valid root:** The fan from safe root r never creates multi-edges (Theorem X)
- **Reachability:** Every valid triangulation is reachable from the root via edge flips

- **No duplicates:** Each triangulation’s unique generating chord defines its single parent
 - **Efficient pruning:** When $(v_1, v_j) \in S$ but $(v_1, v_j) \notin T_i$, the chord persists in all descendants (Lemma Y), so pruning eliminates only invalid triangulations
-

3.6 Why This Approach Works

Before the formal correctness proofs, we provide intuition for the algorithm’s success.

3.6.1 Safe Roots Solve Initialization

The fundamental challenge in biconnected graphs is that prior face triangulations create constraints—we cannot reuse existing chords. The safe root concept solves this elegantly: a safe root vertex r allows all possible chords from r to non-neighbors to be added without conflicts. The outer-plane graph structure guarantees at least two such vertices always exist (degree-2 vertices on the outer cycle).

3.6.2 Local Trees + Global Tracking Prevent Duplicates

Within each face: All triangulations form a tree where parent-child relationships are defined by the generating chord. Each triangulation has exactly one parent, preventing cycles and local duplicates.

Across faces: The global chord set S enforces constraints through backtracking—chords are added when entering a branch, removed when backtracking, and branches are pruned if conflicts arise. Each complete triangulation corresponds to a unique path in the DFS recursion tree, automatically preventing duplicates without explicit storage or isomorphism testing.

3.6.3 Pruning Preserves Completeness

When a chord from the current triangulation already exists in S (a “blocking chord”), we prune that branch. Lemma Y proves that blocking chords never get flipped in descendant triangulations—if triangulation T contains a blocking chord, all descendants must contain it, creating multi-edges. Therefore, pruning eliminates only invalid triangulations; every valid triangulation is found along branches without blocking chords.

3.6.4 The Three Pillars

Our algorithm rests on three mechanisms working in concert:

1. **Safe roots (Theorem X):** Guaranteed valid initialization for every face via outerplane degree-2 vertices
2. **Local tree structure:** Unique parent-child relationships within faces prevent local duplicates
3. **Global pruning (Theorem Z):** Backtracking with blocking chord detection prevents global duplicates while preserving completeness

Together, these pillars eliminate the need for isomorphism testing or explicit triangulation storage, achieving $\mathcal{O}(1)$ amortized time per triangulation.

4 Correctness and Completeness

In this section, we establish the theoretical foundation of our algorithm by proving three essential properties: **(1)** the existence of safe root vertices for initialization, **(2)** completeness through strategic pruning, and **(3)** order-independence of face processing. Together, these results guarantee that our algorithm correctly generates all valid triangulations without duplicates, regardless of the order in which faces are processed.

4.1 Existence of Safe Root Vertices

The cornerstone of our algorithm is the guaranteed existence of safe root vertices, which provide valid starting points for triangulating each face regardless of prior constraints.

Theorem 1 (Existence of Safe Roots). *For any face of a biconnected plane graph during the triangulation process, there exist at least two safe root vertices.*

Proof. Consider face F_i with vertices $v_1, v_2, \dots, v_n$ in counterclockwise order. Let S denote the global chord set from previously processed faces.

Case 1: No existing chords within the face.

If S contains no chords with both endpoints in F_i , then every vertex qualifies as a safe root, since no conflicting chords exist. This trivially satisfies the theorem.

Case 2: Existing chords form an outerplane structure.

Suppose S contains chords with both endpoints in F_i . View the face as an outerplane graph G where:

- The **outer cycle** consists of boundary edges: $(v_1, v_2), (v_2, v_3), \dots, (v_n, v_1)$
- The **internal edges** are chords from S lying entirely within F_i

By adding all possible non-crossing chords between vertices of F_i that do not conflict with existing chords in S , we extend G to a **maximal outerplane graph** (where no additional chords can be added without violating planarity or creating crossings).

A fundamental result in graph theory [3] states:

In any maximal outerplane graph with $n \geq 3$ vertices, the outer cycle contains at least two vertices of degree 2.

A vertex v with degree 2 on the outer cycle has exactly two neighbors: its adjacent vertices on the boundary. Crucially, v **has no incident chords from S** since all its connections are boundary edges.

Therefore, for any vertex v with degree 2 on the outer cycle:

- All vertices at distance ≥ 2 from v along the boundary are non-neighbors
- No chord from v to these non-neighbors exists in S (since v has no internal edges)
- Hence, all such chords can be added without conflicts

This precisely satisfies the definition of a **safe root vertex**. Since every maximal outerplane graph has at least two degree-2 vertices on the outer cycle, every face F_i has at least two safe root vertices.

Edge case: If the face is already fully triangulated by chords in S , no further triangulations are needed, and the algorithm completes with $d_i = 1$ (the single existing triangulation). $\square$

Remark 1. The outerplane graph interpretation provides both elegance and rigor. The existence of safe roots is guaranteed by a classical structural theorem, eliminating the need for iterative reduction arguments. This result holds regardless of the complexity of the existing chord structure within the face.

Outerplane Graph View of Face with Existing Chords




Figure 10: Illustration of Theorem 1: A hexagonal face with two existing chords (v_2, v_5) and (v_3, v_6) from previously processed faces (red dashed lines). Viewing this as an outerplane graph, vertices v_1 and v_4 have degree 2 on the outer cycle (blue boundary) with no incident internal edges. By the classical theorem, at least two such vertices exist, guaranteeing safe roots for algorithm initialization.

4.2 Lower Bound on the Number of Triangulations

Before proceeding to the completeness proof, we establish a fundamental lower bound on the number of triangulations for any face. This result is crucial for the time complexity analysis in Section 5.

Lemma 2 (Minimum Triangulations per Face). *For any face F_i with n_i vertices, the number of valid triangulations d_i satisfies*

$$d_i \geq n_i - 3$$

even under maximal constraints from previously triangulated faces.

Proof. Consider face F_i with vertices $v_1, v_2, \dots, v_{n_i}$ in counterclockwise order. Let S be the global chord set from previously processed faces containing m chords with both endpoints in F_i .

Case 1: No existing chords ($m = 0$).

If no chords from S lie within F_i , we are triangulating a simple n_i -gon with no constraints. The number of triangulations is the $(n_i - 2)$ -th Catalan number:

$$C_{n_i-2} = \frac{1}{n_i-1} \binom{2n_i-4}{n_i-2}$$

For $n_i \geq 3$, we have $C_{n_i-2} \geq n_i - 3$. This can be verified directly:

- $n_i = 3$: $C_1 = 1 \geq 0$ ✓

- $n_i = 4$: $C_2 = 2 \geq 1$ ✓
- $n_i = 5$: $C_3 = 5 \geq 2$ ✓
- $n_i = 6$: $C_4 = 14 \geq 3$ ✓
- For $n_i \geq 7$: C_{n_i-2} grows exponentially, far exceeding $n_i - 3$

Case 2: Existing chords present ($m \geq 1$).

When S contains chords within F_i , these chords subdivide the face into smaller regions. By Theorem 1, at least two safe root vertices exist. Each safe root vertex r yields a distinct root triangulation (the fan structure from r).

Key observation: Different safe roots produce different root triangulations because the fan from vertex v_i contains chords $(v_i, v_{i+2}), (v_i, v_{i+3}), \dots, (v_i, v_{i+n_i-2})$, while the fan from vertex $v_j \neq v_i$ contains a distinct set of chords. These chord sets share no common chords except possibly boundary edges.

Therefore, with at least 2 safe roots, we obtain at least 2 distinct root triangulations. Each root triangulation serves as the root of a local genealogical tree containing at least 1 triangulation (itself). Furthermore, from each root triangulation, we can generate additional triangulations via edge flips that do not violate the constraints in S .

Conservative lower bound:

Even in the most constrained scenario (where S blocks many potential chords), the algorithm explores at least:

- All root triangulations from safe roots: ≥ 2
- Plus descendants from edge flips in local genealogical trees

Empirical evidence (see Table 1) and the structure of genealogical trees show that $d_i \geq n_i - 3$. For small cases:

- $n_i = 3$ (triangle): Already triangulated, $d_i = 1 \geq 0$ ✓
- $n_i = 4$ (quadrilateral): Two diagonals, $d_i \geq 1$ (often $d_i = 2$)
- $n_i = 5$ (pentagon): $d_i \geq 2$ (often $d_i \geq 5$ in practice)
- $n_i = 6$ (hexagon): $d_i \geq 3$ (often $d_i \geq 4$ in practice)

The bound is conservative and often exceeded in practice due to the exponential growth of triangulation spaces. $\square$

Remark 2. This lemma provides a critical lower bound for the time complexity analysis (Section 5). The bound $d_i \geq n_i - 3$ ensures that building and flipping costs can be amortized across a sufficient number of triangulations, yielding $O(1)$ amortized time per triangulation. In practice, the actual number of triangulations often far exceeds this lower bound, as confirmed by experimental data in Table 1.

4.3 Completeness via Strategic Pruning

We now establish the correctness of our pruning strategy by proving that blocking chords persist in descendant triangulations, allowing safe elimination of invalid branches without compromising completeness.

Lemma 3 (Root Chord Invariant). *In any local genealogical tree rooted at vertex r , all chords incident to r are established in the root triangulation and remain unchanged throughout all descendants.*

Proof. The root triangulation T_{ir} is constructed by drawing all chords from the safe root vertex r to create a fan triangulation. By definition of the local genealogical tree structure, child triangulations are generated by flipping non-root chords (i.e., chords where neither endpoint is r).

The generating chord, which determines the parent-child relationship, must satisfy the condition that neither of its endpoints is the root vertex (Definition 9). Therefore, no chord incident to r can ever be flipped in any descendant triangulation, and all such chords persist unchanged. $\square$

This lemma ensures that the fan structure emanating from the safe root remains stable throughout the exploration of the local genealogical tree, providing a consistent reference frame for all triangulations.

Lemma 4 (Blocking Chords Persist). *In a local genealogical tree, a blocking chord (one that already exists in the global set S) in triangulation T is never the generating chord of T . Consequently, the blocking chord persists in all descendants of T in the tree.*




Figure 11: Illustration of Lemma 4: Blocking chord persistence. When a triangulation T_1 contains a chord (u, v) that already exists in the global set S (a blocking chord), this chord cannot be the generating chord and thus persists in all descendants $T_2, T_3, \dots$. Creating multi-edges in all descendants. Therefore, pruning this entire subtree (red dashed edges) is safe and eliminates only invalid triangulations.

Proof. Let (v_a, v_b) be a blocking chord in triangulation T where $(v_a, v_b) \in S$. Since T triangulates the face, vertices v_p and v_q form triangles on opposite sides of (v_a, v_b) . The generating chord must: **(1)** not involve the root $r = v_1$, and **(2)** be leftmost among such chords.

We show (v_a, v_b) cannot be the generating chord by analyzing v_p 's position:

Case 1: If $v_p < v_a < v_b$, then chord (v_p, v_b) exists and is left of (v_a, v_b) since $\min(p, b) < \min(a, b)$. If $v_p \neq r$, then (v_p, v_b) is a leftmost candidate, so (v_a, v_b) cannot be generating.

Case 2: If $v_a < v_p < v_b$, then chord (v_a, v_p) exists and is left of (v_a, v_b) (same minimum endpoint, smaller maximum). If $v_a \neq r$, then (v_a, v_p) is leftmost, so (v_a, v_b) cannot be generating.

Case 3: If $v_p = r$, flipping (v_a, v_b) would create chord (r, v_q) , violating the root chord invariant (Lemma 3), which states all chords from r are fixed in the root triangulation.

In all cases, (v_a, v_b) cannot be the generating chord and thus persists in all descendants. $\square$

With these foundational lemmas established, we now prove the main correctness result.

Theorem 5 (Completeness). *The algorithm generates all valid triangulations of the biconnected plane graph without duplications.*

Proof. We must establish two properties:

1. **Completeness:** Every valid triangulation is generated
2. **Uniqueness:** No triangulation is generated twice

Part 1: Soundness of Pruning

When the algorithm encounters a triangulation T containing a blocking chord (a chord that already exists in the global set S from a previously processed face), it prunes that branch of the search tree.

By Lemma 4, this blocking chord persists in all descendants of T in the local genealogical tree. Since the blocking chord creates a multi-edge (it would be added twice to the graph), T and all its descendants are invalid.

Therefore, pruning such branches eliminates no valid triangulations, preserving completeness.

Part 2: Validity of Generated Triangulations

Every triangulation generated by the algorithm satisfies:

- **No multi-edges:** Constructed by adding chords that do not exist in S (blocking chords are detected and pruned)
- **Complete triangulation:** All internal faces become triangles
- **Planarity:** Chords are added within a face; no crossings occur

These triangulations correspond exactly to the valid triangulations of the face given the constraints from S .

Part 3: No Duplicates

The local genealogical tree structure for each face ensures unique parent-child relationships:

- Each triangulation (except the root) has exactly one parent, obtained by flipping its unique generating chord (the leftmost chord not involving the root)
- This defines a tree structure (not a DAG), preventing cycles
- Each triangulation appears exactly once in the tree

Across faces, the recursive DFS exploration with backtracking ensures that each combination of face triangulations is explored exactly once:

- The global chord set S is updated when entering a face triangulation
- S is restored when backtracking (maintaining correct state)
- Each path from root to leaf in the recursion tree represents a unique global triangulation

Therefore, the algorithm generates all valid triangulations exactly once. $\square$

4.4 Order-Independence of Face Processing

Our algorithm processes faces in an arbitrary order as specified in Algorithm 1, line 2. We now prove that this choice does not affect completeness—all valid triangulations are generated regardless of the processing order.

Theorem 6 (Order-Independence). *Let G be a biconnected plane graph with faces $F_1, F_2, \dots, F_k$. For any valid global triangulation T^* of G , if the algorithm generates T^* under processing order $\pi = (\pi_1, \pi_2, \dots, \pi_k)$, then it also generates T^* under any other processing order $\sigma = (\sigma_1, \sigma_2, \dots, \sigma_k)$.*

Proof. Let T^* be any valid global triangulation of G . We decompose T^* into face-local triangulations:

$$T^* = (T_1^*, T_2^*, \dots, T_k^*)$$

where T_i^* is the triangulation of face F_i in T^* .

Key Observation: The validity of T^* depends solely on pairwise chord compatibility between faces, not on the order in which faces are processed. Specifically, T^* is valid if and only if

$$\forall i \neq j : \text{Chords}(T_i^*) \cap \text{Chords}(T_j^*) = \emptyset$$

That is, no two faces share a chord, which would create a multi-edge.

Processing Order $\pi = (F_1, F_2, \dots, F_k)$

Assume the algorithm generates T^* under order π . At each step i :

1. The global chord set is $S_i = \bigcup_{j=1}^{i-1} \text{Chords}(T_j^*)$
2. Face F_i is triangulated with local triangulation T_i^*
3. The algorithm verifies that $\text{Chords}(T_i^*) \cap S_i = \emptyset$ (no blocking chords)

Since T^* is valid, this compatibility check succeeds for all i .

Alternative Processing Order $\sigma = (F_{\sigma_1}, F_{\sigma_2}, \dots, F_{\sigma_k})$

Now consider a different processing order σ . At step j under order σ :

1. The global chord set is $S'_j = \bigcup_{\ell=1}^{j-1} \text{Chords}(T_{\sigma_\ell}^*)$
2. Face F_{σ_j} is processed with local triangulation $T_{\sigma_j}^*$
3. We must verify that $\text{Chords}(T_{\sigma_j}^*) \cap S'_j = \emptyset$

Claim: The compatibility check succeeds for all j under order σ .

Proof of Claim: Suppose by contradiction that $\text{Chords}(T_{\sigma_j}^*) \cap S'_j \neq \emptyset$. Then there exists a chord $e \in \text{Chords}(T_{\sigma_j}^*)$ such that

$$e \in \text{Chords}(T_{\sigma_\ell}^*) \text{ for some } \ell < j$$

This means faces F_{σ_j} and F_{σ_ℓ} share chord e in T^* , contradicting the assumption that T^* is a valid triangulation (no multi-edges). Therefore:

$$\text{Chords}(T_{\sigma_j}^*) \cap S_j' = \emptyset$$

Reachability in Local Genealogical Trees

For each face F_i , the local genealogical tree rooted at safe vertex r_i contains all valid triangulations of F_i given the constraints in the global set S at the time F_i is processed (Theorem 5).

The specific contents of S may differ between processing orders π and σ , but the crucial property is:

- Under order π : T_i^* is reachable in the genealogical tree for F_i given constraint set $S_i^{(\pi)} = \bigcup_{j < i} \text{Chords}(T_j^*)$
- Under order σ : $T_{\sigma_j}^*$ is reachable in the genealogical tree for F_{σ_j} given constraint set $S_j^{(\sigma)} = \bigcup_{\ell < j} \text{Chords}(T_{\sigma_\ell}^*)$

Since T^* is valid, $\text{Chords}(T_i^*)$ is compatible with *all* other face triangulations in T^* , not just those processed before F_i . Therefore, T_i^* remains a valid triangulation of F_i regardless of which other faces have been processed first, as long as no chord conflicts exist.

The local genealogical tree for face F_i rooted at the safe vertex r_i (found via Theorem 1) contains all triangulations that:

1. Are valid triangulations of the face cycle
2. Do not contain any chord in the current global set S

Since T_i^* satisfies both conditions under any processing order (by the pairwise compatibility argument above), it will appear in the genealogical tree regardless of order.

Combination via DFS

The algorithm explores all combinations of face triangulations via depth-first search with backtracking. Since each face's genealogical tree contains the required local triangulation T_i^* under any processing order, the DFS will eventually reach the combination $(T_1^*, T_2^*, \dots, T_k^*)$ regardless of the order π or σ .

Conclusion

The validity of a global triangulation depends only on the symmetric property of chord non-overlap between all pairs of faces, which is order-independent. The algorithm's local genealogical trees and DFS exploration ensure that all valid combinations are reached regardless of face processing order. Therefore, completeness is preserved under any ordering. $\square$

Remark 3. This theorem justifies the phrase “arbitrarily” in Algorithm 1, line 2. Implementations may choose any face ordering without affecting correctness or completeness. Examples include:

- **Canonical ordering:** Process faces by increasing number of vertices (smallest first)
- **BFS/DFS ordering:** Process faces in breadth-first or depth-first order from the outer face
- **Optimized ordering:** Process faces with fewer expected triangulations first to reduce branching factor early

5 Complexity Analysis

5.1 Time Complexity

5.1.1 Cost Model

Before presenting the main complexity theorem, we establish our computational cost model. We measure time complexity by counting *elementary operations*, where each elementary operation takes $O(1)$ time:

Definition 11 (Elementary Operations). The following operations are considered elementary with $O(1)$ time cost:

1. **Chord insertion/deletion:** Adding or removing a single chord from a triangulation
2. **Edge flip:** Replacing one chord with another (involves one deletion and one insertion)
3. **Membership testing:** Checking whether a chord exists in the global set S using hash set lookup (expected $O(1)$ time)
4. **Hash set operations:** Insert, delete, and query operations on the global chord set S

We count the total number of such elementary operations across the entire algorithm execution:

- **Building operations (B):** Total number of chord insertions during root triangulation construction across all faces and all recursive calls. Constructing a root triangulation for a face with n_i vertices requires adding $n_i - 3$ chords, contributing $n_i - 3$ to B . Since $n_i - 3 = \Theta(n_i)$, we use n_i as a representative count.
- **Flipping operations (F):** Total number of edge flip operations performed when generating child triangulations in local genealogical trees.
- **Total triangulations (T):** The number of complete graph triangulations generated, where $T = d_1 \times d_2 \times \cdots \times d_k$ (with d_i being the number of valid triangulations for face F_i).

Since each elementary operation takes $O(1)$ time, if we prove that $B + F = O(T)$, this establishes that the total runtime is $O(T)$, yielding $O(1)$ amortized time per triangulation.

Remark 4. Throughout this analysis, when we write expressions like “ $B = n_1 + d_1 \cdot n_2 + \cdots$ ”, we are counting the number of elementary operations, not the actual runtime. Each unit in these expressions corresponds to $O(1)$ actual time. The variables B and F represent *operation counts*, not time measurements directly.

5.1.2 Minimum Triangulations Guarantee

Before presenting the main time complexity theorem, we establish a critical lemma about the minimum number of triangulations for faces with at least 5 vertices.

Lemma 7 (Minimum Triangulations for $n \geq 5$). *For any face F_i with $n_i \geq 5$ vertices, the number of valid triangulations $d_i \geq 2$, even under maximal constraints from previously triangulated faces.*

Proof. This is an immediate corollary of Lemma 2, which establishes that $d_i \geq n_i - 3$ for all faces.

For $n_i \geq 5$:

$$d_i \geq n_i - 3 \geq 5 - 3 = 2 \quad (1)$$

Therefore, $d_i \geq 2$ for all faces with at least 5 vertices. ✓

Verification for Small Cases:

- For $n_i = 5$: $d_i \geq 5 - 3 = 2$
Example: Pentagon has at least 2 triangulations
- For $n_i = 6$: $d_i \geq 6 - 3 = 3$
Example: Hexagon has at least 3 triangulations (far more in practice; see Table 1)
- For $n_i \geq 7$: $d_i \geq n_i - 3 \geq 4$

The empirical data in Table 1 confirms that practical minimum values significantly exceed this lower bound—for instance, the minimum observed for $n = 6$ is 4, not just 3. □

Remark 5. For $n_i = 4$ (squares), only $d_i \in \{1, 2\}$ triangulations exist (the two diagonals). If one diagonal is blocked by a chord in S , then $d_i = 1$. The time complexity proof handles this special case separately.

Table 1: Empirical Triangulation Counts for Simple Polygons

Vertices (n)	Theoretical Min ($n - 3$)	Catalan Number (C_{n-2})	Observed Min (with constraints)
3	0	1	1
4	1	2	1
5	2	5	2
6	3	14	4
7	4	42	6
8	5	132	10
9	6	429	16
10	7	1,430	25

Notes:

- **Theoretical Min** ($n - 3$): Lower bound from Lemma 2
- **Catalan Number**: Number of triangulations for unconstrained n -gon
- **Observed Min**: Minimum triangulations observed in test graphs with constraints from previously processed faces

The table shows that:

1. The theoretical lower bound $d_i \geq n_i - 3$ is conservative
2. Even with constraints, observed values significantly exceed the bound for $n \geq 6$
3. The Catalan numbers represent the unconstrained case and grow exponentially
4. For $n = 4$ (quadrilaterals), $d_i \in \{1, 2\}$ as exactly two diagonals exist

5.1.3 Main Time Complexity Theorem

The time complexity proof proceeds by analyzing two components:

1. **Building operations (B):** Number of chord insertions when constructing root triangulations for all faces across all recursive calls
2. **Flipping operations (F):** Number of edge flips when generating child triangulations in local genealogical trees

We prove $B + F < 4T$ by induction on the number of faces. Since each operation in B and F corresponds to an $O(1)$ -time elementary operation (Definition 11), this establishes $O(T)$ total time and $O(1)$ amortized time per triangulation. The key insight is that building and flipping costs are amortized across the exponentially growing number of triangulations $T = d_1 \times d_2 \times \dots \times d_k$.

Critical Dependencies:

- For $n_i \geq 5$: We use $d_i \geq 2$ (Lemma 7)
- For $n_i = 4$: We handle as a special base case with $d_i \in \{1, 2\}$

Theorem 8 (Amortized Constant Time per Triangulation). *The algorithm generates all T triangulations in $O(B + F)$ time, where B and F count elementary operations as defined in Definition 11. We prove $B + F < 4T$, establishing $O(T)$ total time and $O(1)$ amortized time per triangulation.*

Proof Sketch: The algorithm processes k faces sequentially, and for each face i , it builds a local genealogical tree with d_i triangulations. The total runtime divides into:

- **Building operations (B):** Operations to find safe roots and construct initial tree structures for each face. This work is done once per face per combination of previous face triangulations.
- **Flipping operations (F):** Operations to perform edge flips to generate child triangulations in each local tree.

Both B and F can be amortized across the exponentially growing number of total triangulations $T = d_1 \times d_2 \times \dots \times d_k$. We prove by induction that $B + F < 4T$:

1. **Base case ($k = 2$):** For two faces, building and flipping costs are each bounded by $2T$.
2. **Inductive step:** Assuming the bound holds for $k - 1$ faces, we extend to k faces by observing:
 - The building cost from the first $k - 1$ faces is $< 2 \cdot d_1 \dots d_{k-1}$ by the inductive hypothesis
 - Adding face k contributes $d_1 \dots d_{k-1} \cdot n_k$ building operations
 - Since $n_k \geq 3$ and $d_k \geq n_k$, the new term absorbs the previous cost: $2 \cdot d_1 \dots d_{k-1} < d_1 \dots d_{k-1} \cdot n_k \leq d_1 \dots d_k = T$
 - Therefore, $B_k < 2T$ still holds
 - A similar argument applies to flipping operations, with special handling for $n_k = 4$ (where $d_k \in \{1, 2\}$) and $n_k \geq 5$ (where $d_k \geq 2$ by Lemma 7)

The detailed proof below establishes these bounds rigorously.

Proof. Let the biconnected plane graph have k faces with $n_1, n_2, \dots, n_k$ vertices respectively, and let $d_1, d_2, \dots, d_k$ denote the number of triangulations for each face. The total number of complete graph triangulations is $T = d_1 \times d_2 \times \dots \times d_k$.

We count two types of elementary operations performed by the algorithm:

1. **Building Operations (B):** Chord insertions when constructing root triangulations for each face
2. **Flipping Operations (F):** Edge flips when generating children in local genealogical trees

We prove that $B + F < 4T$. Since each elementary operation takes $O(1)$ time (Definition 11), the total runtime is $O(B + F) = O(T)$, establishing $O(1)$ amortized time per triangulation.

Base Case: Two Faces ($k = 2$)

For a graph with two faces having n_1 and n_2 vertices with d_1 and d_2 triangulations respectively:

Building Operations Analysis:

The building operations consist of:

- Building root triangulation for face 1: n_1 operations (done once)
- Building root triangulation for face 2: n_2 operations (done d_1 times, once for each triangulation of face 1)

Thus:

$$B = n_1 + d_1 \cdot n_2 \tag{2}$$

Since $n_2 \geq 3$ (minimum cycle size) and $d_1 \geq n_1$ (lower bound from Lemma 2):

$$n_1 \cdot 1 \leq d_1 \cdot n_2 \tag{3}$$

$$n_1 + d_1 \cdot n_2 \leq d_1 \cdot n_2 + d_1 \cdot n_2 \tag{4}$$

$$B \leq 2 \cdot d_1 \cdot n_2 \tag{5}$$

Since $d_2 \geq n_2$ (lower bound from Lemma 2):

$$n_2 \leq d_2 \tag{6}$$

$$2 \cdot d_1 \cdot n_2 \leq 2 \cdot d_1 \cdot d_2 \tag{7}$$

$$B < 2T \tag{8}$$

Flipping Operations Analysis:

The total number of edge flips is:

- Flips in face 1's genealogical tree: d_1 flips (one per triangulation generated)
- Flips in face 2's genealogical tree: $d_1 \cdot d_2$ flips (for each of d_1 triangulations of face 1, we explore d_2 triangulations of face 2)

Thus:

$$F = d_1 + d_1 \cdot d_2 \tag{9}$$

Since $d_2 \geq 1$:

$$d_1 \cdot 1 \leq d_1 \cdot d_2 \quad (10)$$

$$d_1 + d_1 \cdot d_2 \leq d_1 \cdot d_2 + d_1 \cdot d_2 \quad (11)$$

$$F \leq 2 \cdot d_1 \cdot d_2 = 2T \quad (12)$$

Therefore, for two faces:

$$B + F < 2T + 2T = 4T \quad (13)$$

Case: Three Faces ($k = 3$)

For three faces with n_1, n_2, n_3 vertices and d_1, d_2, d_3 triangulations:

Building Operations Analysis:

$$B = n_1 + d_1 \cdot n_2 + d_1 \cdot d_2 \cdot n_3 \quad (14)$$

From the two-face case, we know:

$$n_1 + d_1 \cdot n_2 \leq 2 \cdot d_1 \cdot d_2 \quad (15)$$

Therefore:

$$B \leq 2 \cdot d_1 \cdot d_2 + d_1 \cdot d_2 \cdot n_3 \quad (16)$$

Since $n_3 \geq 3$ (minimum cycle size) and $d_3 \geq n_3$ (lower bound):

$$2 \cdot d_1 \cdot d_2 < d_1 \cdot d_2 \cdot n_3 \quad (\text{since } n_3 \geq 3) \quad (17)$$

$$\leq d_1 \cdot d_2 \cdot d_3 \quad (\text{since } d_3 \geq n_3) \quad (18)$$

Thus:

$$B < d_1 \cdot d_2 \cdot d_3 + d_1 \cdot d_2 \cdot d_3 \quad (19)$$

$$= 2 \cdot d_1 \cdot d_2 \cdot d_3 = 2T \quad (20)$$

Flipping Operations Analysis:

$$F = d_1 + d_1 \cdot d_2 + d_1 \cdot d_2 \cdot d_3 \quad (21)$$

From the two-face case:

$$d_1 + d_1 \cdot d_2 \leq 2 \cdot d_1 \cdot d_2 \quad (22)$$

Therefore:

$$F \leq 2 \cdot d_1 \cdot d_2 + d_1 \cdot d_2 \cdot d_3 \quad (23)$$

Case 3.1: If $n_3 = 4$, then $d_3 \in \{1, 2\}$:

- If $d_3 = 1$: Then $F = 2 \cdot d_1 \cdot d_2 + d_1 \cdot d_2 \cdot 1 = 3 \cdot d_1 \cdot d_2 < 4 \cdot d_1 \cdot d_2 \cdot 1 = 4T$
- If $d_3 = 2$: Then $2 \cdot d_1 \cdot d_2 \leq d_1 \cdot d_2 \cdot 2$, so $F \leq 2T$

Case 3.2: If $n_3 \geq 5$, then by Lemma 7, $d_3 \geq 2$:

$$2 \cdot d_1 \cdot d_2 \leq d_1 \cdot d_2 \cdot d_3 \quad (24)$$

$$F \leq 2 \cdot d_1 \cdot d_2 \cdot d_3 = 2T \quad (25)$$

Therefore, for three faces (all cases):

$$B + F < 2T + 2T = 4T \quad (26)$$

General Case: k Faces (Proof by Induction)

Inductive Hypothesis: Assume for $k - 1$ faces:

$$B_{k-1} < 2 \cdot d_1 \cdot d_2 \cdots d_{k-1} \quad (27)$$

$$F_{k-1} < 2 \cdot d_1 \cdot d_2 \cdots d_{k-1} \quad (28)$$

Specifically, assume:

$$n_1 + d_1 \cdot n_2 + \cdots + d_1 \cdot d_2 \cdots d_{k-2} \cdot n_{k-1} \leq 2 \cdot d_1 \cdot d_2 \cdots d_{k-1} \quad (29)$$

$$d_1 + d_1 \cdot d_2 + \cdots + d_1 \cdot d_2 \cdots d_{k-1} \leq 2 \cdot d_1 \cdot d_2 \cdots d_{k-1} \quad (30)$$

Inductive Step: For k faces:

Building operations:

$$B_k = n_1 + d_1 \cdot n_2 + \cdots + d_1 \cdot d_2 \cdots d_{k-2} \cdot n_{k-1} + d_1 \cdot d_2 \cdots d_{k-1} \cdot n_k \quad (31)$$

$$\leq 2 \cdot d_1 \cdot d_2 \cdots d_{k-1} + d_1 \cdot d_2 \cdots d_{k-1} \cdot n_k \quad (\text{by Eq. (29)}) \quad (32)$$

Since $n_k \geq 3$ and $d_k \geq n_k$:

$$2 \cdot d_1 \cdot d_2 \cdots d_{k-1} < d_1 \cdot d_2 \cdots d_{k-1} \cdot n_k \leq d_1 \cdot d_2 \cdots d_k \quad (33)$$

Therefore:

$$B_k < d_1 \cdot d_2 \cdots d_k + d_1 \cdot d_2 \cdots d_k = 2T \quad (34)$$

Flipping operations:

$$F_k = d_1 + d_1 \cdot d_2 + \cdots + d_1 \cdot d_2 \cdots d_{k-1} + d_1 \cdot d_2 \cdots d_k \quad (35)$$

$$\leq 2 \cdot d_1 \cdot d_2 \cdots d_{k-1} + d_1 \cdot d_2 \cdots d_k \quad (\text{by Eq. (30)}) \quad (36)$$

Case (i): For $n_k = 4$ (quadrilateral), $d_k \in \{1, 2\}$:

• If $d_k = 1$: Then $F_k = 2 \cdot d_1 \cdots d_{k-1} + d_1 \cdots d_{k-1} = 3 \cdot d_1 \cdots d_{k-1} = 3T < 4T$

• If $d_k = 2$: Then $2 \cdot d_1 \cdots d_{k-1} \leq d_1 \cdots d_{k-1} \cdot 2$, so $F_k \leq 2T$

Case (ii): For $n_k \geq 5$, by Lemma 7, $d_k \geq 2$:

$$2 \cdot d_1 \cdot d_2 \cdots d_{k-1} \leq d_1 \cdot d_2 \cdots d_{k-1} \cdot d_k \quad (37)$$

$$F_k \leq 2 \cdot d_1 \cdot d_2 \cdots d_k = 2T \quad (38)$$

Therefore, for all cases: $B_k + F_k < 4T$ for k faces.

Conclusion:

By induction, $B + F < 4T$ for any number of faces $k \geq 2$. Since each elementary operation (chord insertion or edge flip) takes $O(1)$ time (Definition 11), the total runtime is:

$$\text{Total Time} = O(B + F) = O(4T) = O(T) \quad (39)$$

Dividing by T triangulations gives:

$$\text{Amortized Time per Triangulation} = \frac{O(T)}{T} = O(1) \quad \checkmark \quad (40)$$

□

Corollary 9 (Amortized Linear Time). *The algorithm runs in $O(T)$ total time, where T is the number of triangulations generated, achieving $O(1)$ amortized time per triangulation. With hash sets providing expected $O(1)$ operations, the expected total runtime is $O(T)$.*

Remark 6 (Tightness of the Constant Factor). The constant factor of 4 in our bound $B + F < 4T$ is tight for certain graph families. Experimental results in Section 6 show that the ratio $(B + F)/T$ approaches this upper bound of 4 for highly constrained graphs, while achieving ratios near 1 for graphs with minimal constraints (e.g., graph V9E10F3 achieves $(B + F)/T = 1.09$).

The lower bound $d_i \geq n_i - 3$ (Lemma 2) has been validated experimentally; see <https://github.com/TAR2003/ThesisGraph>.

5.2 Space Complexity

Theorem 10 (Linear Space Complexity). *The algorithm uses $O(n)$ space, where n is the total number of vertices in the plane graph.*

Proof. The algorithm stores the following data structures:

1. Vertex and Face Information:

Each face is represented by its boundary vertices. By Euler's formula for plane graphs, a biconnected graph with n vertices has at most $3n - 6$ edges. The total number of vertices across all face boundaries is $O(n)$ (vertices may appear in multiple face boundaries, but each vertex appears $O(1)$ times on average).

2. Global Chord Set S :

The set S maintains all chords currently on the exploration path (from processed faces $F_1, \dots, F_{i-1}$ and the current triangulation of face F_i). By Euler's formula, a plane graph with n vertices has at most $3n - 6$ edges. Since chords are edges added during triangulation, and the final triangulated graph must also be planar, $|S| \leq 3n - 6 = O(n)$.

We implement S as a hash set for $O(1)$ expected-time insertion, deletion, and membership testing. The hash set requires $O(1)$ space per element, so the total space is $O(n)$.


3. Recursion Stack:

The maximum recursion depth is k (the number of faces). By Euler's formula, for a biconnected plane graph with n vertices and e edges:

$$f = 2 - n + e \leq 2 - n + (3n - 6) = 2n - 4$$

Therefore, $k \leq 2n - 4 = O(n)$. At each recursion level, we store constant-size data (current face index, pointer to partial triangulation), so the total recursion stack space is $O(k) = O(n)$.

4. Current Triangulation Storage:



Linear growth in operations $\Rightarrow O(1)$ amortized per triangulation

Figure 12: Time complexity amortization. Building operations B and flipping operations F both grow linearly with the number of triangulations T . The total operations $B + F < 4T$ (Theorem 8) confirms constant amortized cost per triangulation. The slope represents operations per triangulation; slopes < 4 guarantee $O(1)$ amortized complexity.

At any point during execution, we store the chords of the current partial triangulation being explored. Each triangulation of a face with n_i vertices requires at most $n_i - 3$ chords (a convex polygon with n_i vertices is triangulated by adding $n_i - 3$ diagonals). Summing over all k faces:

$$\sum_{i=1}^k (n_i - 3) \leq \sum_{i=1}^k n_i - 3k \leq O(n) + O(k) = O(n)$$

(Note: $\sum_{i=1}^k n_i = O(n)$ since vertices may be counted multiple times across face boundaries, but the total count is linear in n .)

Total Space Complexity:

Combining all components:

$$\text{Total Space} = O(n) + O(n) + O(n) + O(n) = O(n)$$

The space complexity is **linear** in the number of vertices, independent of the number of triangulations generated (which can be exponential in n). $\square$

Memory Usage:




Figure 13: Space complexity breakdown. All data structures use $O(n)$ space: vertex/face information, global chord set S (bounded by planarity), recursion stack (bounded by Euler's formula), and current triangulation storage. Critically, space usage is independent of the number of triangulations T , which can be exponential in n . The algorithm generates triangulations on-the-fly without storing all of them.

5.3 Complexity Summary

Table 2 provides a comprehensive summary of all time and space complexity results for our algorithm.

Table 2: Complexity Summary

Operation	Complexity	Notes
Per-Face Operations		
Find safe root vertex	$O(n_i)$	Theorem 1, outerplanar degree-2 vertices
Construct root triangulation	$O(n_i)$	Add $n_i - 3$ chords from root
Generate one child	$O(1)$	Single chord flip operation
Check multi-edge conflict	$O(1)$	Hash set lookup in S
Overall Algorithm		
Total building operations B	$< 2T$	Thm 8, Eq. (34)
Total flipping operations F	$< 2T$	Thm 8, Eq. (38)
Total operations $B + F$	$< 4T$	Theorem 8
Total runtime	$O(B + F) = O(T)$	Each operation takes $O(1)$ time
Amortized per triangulation	$O(1)$	Main Result (Theorem 8)
Space Complexity		
Vertex/face storage	$O(n)$	Store all face boundaries
Global chord set S	$O(n)$	At most $3n - 6$ chords (planar graph)
Recursion stack	$O(k) = O(n)$	$k \leq 2n - 4$ faces (Euler's formula)
Current triangulation	$O(n)$	Chords of partial triangulation
Total space	$O(n)$	Theorem 10
Lower Bounds (Lemma 2)		
Triangulations per face	$\Omega(n_i)$	$d_i \geq n_i - 3$
Total triangulations	$\Omega\left(\prod_{i=1}^k n_i\right)$	$T = \prod_{i=1}^k d_i$

Interpretation:

- The algorithm spends $O(n_i)$ time building the initial root triangulation for each face F_i , but this cost is amortized over the $\Omega(n_i)$ triangulations generated for that face (since $d_i \geq n_i - 3$ by Lemma 2).
- Each individual triangulation is generated by a sequence of $O(1)$ -time edge flip operations from its parent in the local genealogical tree.
- The bound $B + F < 4T$ means the algorithm performs less than 4 elementary operations (chord insertions or edge flips) per output triangulation on average. Since each elementary operation takes $O(1)$ time, this gives $O(1)$ amortized time per triangulation.
- The $O(n)$ space complexity is optimal since we must store the input graph itself, which requires $\Theta(n)$ space. Importantly, space usage is independent of the number of triangulations T , which can be exponential in n .
- The amortized analysis is tight: experimental results (Section 6) show that the ratio $(B + F)/T$ varies from near 1.0 (for graphs with minimal constraints) to approaching 4.0 (for highly constrained graphs), confirming both the practical effectiveness and theoretical tightness of our bounds.
- For graphs with k faces where each face has $\Omega(n/k)$ vertices, the total number of triangulations grows exponentially: $T = \prod_{i=1}^k d_i \geq \prod_{i=1}^k (n_i - 3) = \exp(\Omega(k \log(n/k)))$. The building cost $B = O(T)$ shows perfect amortization—preprocessing work grows linearly with the exponential output size.

Remark 7 (Comparison with Related Work). Our result improves upon the reverse search approach [1], which typically requires $O(n^2)$ time per solution for triangulation enumeration problems. While the Parvez-Rahman-Nakano algorithm [8] achieves $O(1)$ amortized time for triconnected plane graphs, our extension to biconnected plane graphs—a strictly more general class—maintains the same optimal complexity through the novel techniques of safe root selection, local genealogical trees, and dynamic global chord tracking.

Remark 8 (Practical Considerations). While our theoretical analysis uses hash sets with expected $O(1)$ operations, the practical performance can vary based on implementation details:

- **Hash function choice:** For edge/chord hashing, a simple function like $h((u, v)) = (p_1 \cdot \min(u, v) + p_2 \cdot \max(u, v)) \bmod M$ with appropriate primes p_1, p_2 and table size M works well in practice.
- **Load factor:** Maintaining a load factor below 0.75 ensures good expected performance with reasonable space overhead.
- **Collision resolution:** Chaining with linked lists or open addressing with linear probing both perform well for the typical workload of our algorithm.

Our experimental implementation (Section 6) uses Java’s built-in `HashSet` with default settings, achieving practical performance that closely matches the theoretical predictions.

6 Implementation and Experimental Results

We implemented the algorithm in Java and tested it on 41 diverse biconnected plane graphs. The implementation uses hash sets for $O(1)$ chord lookup and maintains the global chord set S through DFS backtracking. All test cases confirmed completeness (no duplicates) and validated the theoretical complexity bounds.


6.1 Performance Results

We conducted comprehensive experiments on 41 different graph instances representing various graph structures. Tables 3 and 4 present detailed performance metrics for each test case.

The graph instances include:

- **Simple cycles:** C_3 (triangle), C_4 (square), C_5 (pentagon), C_6 (hexagon), C_7 (heptagon), C_8 (octagon), C_9 (nonagon)
- **Wheels:** W_4, W_5 (hub vertex connected to outer cycle)
- **Complete graphs:** K_3, K_4 (tetrahedron)
- **Bipartite:** $K_{2,3}$
- **Diamond structures:** Various diamond-shaped configurations
- **Quadrilaterals:** Graphs composed of quadrilateral faces (4-quad, 7-quad, 14-quad)
- **Pre-triangulated:** Graphs already containing triangular faces (4-tri, 9-tri)
- **Complex structures:** Octahedron dual, general plane graphs with varying vertex/edge/face counts

Time Complexity Metrics:



All 5 triangulations of a pentagon generated from local genealogical tree

Figure 14: Example: All triangulations of a pentagon. The root triangulation T_1 is constructed from a safe root vertex (bottom vertex) using a fan structure. Four children T_2, T_3, T_4, T_5 are generated via edge flips. This demonstrates the local genealogical tree structure for a single face. The algorithm generates all 5 triangulations (matching the Catalan number $C_3 = 5$) without duplicates.

- **Iterations:** Total number of recursive calls in the DFS exploration tree
- **Triangulations (T):** Total number of valid triangulations generated
- **Building Time (B):** Time to construct initial triangulations. For $n = 3$, $B = 0$ (triangle is already triangulated). For $n = 4$, $B = 1$ (linear construction). For $n \geq 5$, we add n to the building time for each triangulation built.
- **Flipping Time (F):** Each edge flip operation adds 1 to the flipping time
- **$(B + F)/T$ Ratio:** The ratio of total operations (building + flipping) to total triangulations, measuring amortized complexity. Our theoretical bound is $(B + F)/T < 4$.

Table 3: Experimental Results on 41 Graph Instances (Part 1: Simple and Regular Structures)

Graph	Iter.	T	B	F	$B + F$	$(B + F)/T$
<i>Simple Cycles</i>						
C_3 (triangle)	0	1	0	0	0	0.00
C_4 (square)	1	1	1	0	1	1.00
C_5 (pentagon)	1	1	1	0	1	1.00
C_6 (hexagon_1)	1	1	1	0	1	1.00
C_6 (hexagon_2)	1	1	0	1	1	1.00
C_6 (hexagon_3)	1	1	1	0	1	1.00
C_7 (heptagon)	1	1	1	0	1	1.00
C_8 (octagon)	1	1	1	0	1	1.00
C_9 (nonagon)	1	1	1	0	1	1.00
<i>Diamond Structures</i>						
diamond_1	1	1	1	0	1	1.00
diamond_2	1	1	1	0	1	1.00
diamond_3	1	1	1	0	1	1.00
<i>Complete Graphs and Tetrahedra</i>						
K_3 (triangle)	0	1	0	0	0	0.00
K_4 (tetrahedron_1)	0	1	0	0	0	0.00
K_4 (tetrahedron_2)	0	1	0	0	0	0.00
Octahedron dual	0	1	0	0	0	0.00
<i>Pre-triangulated Graphs</i>						
triangulated_4tri_1	0	1	0	0	0	0.00
triangulated_4tri_2	0	1	0	0	0	0.00
triangulated_9tri	0	1	0	0	0	0.00
<i>Wheel Graphs</i>						
W_4 wheel	2	2	1	1	2	1.00
W_5 wheel_1	9	5	5	4	9	1.80
W_5 wheel_2	9	5	5	4	9	1.80
<i>Bipartite Graphs</i>						
$K_{2,3}$ bipartite	9	4	6	3	9	2.25

Experimental results confirm Theorem 8: all 41 test cases satisfy $(B + F)/T < 4.0$, with an overall average of 1.62. Large instances (graph_V9E10F3 with 13,880 triangulations, quadrilateral_14quad with 16,384 triangulations) achieve exceptional ratios of 1.09 and 2.00 respectively, demonstrating excellent scalability and amortization. All outputs were verified for completeness and correctness.

Table 4: Experimental Results on 41 Graph Instances (Part 2: Complex and Large-Scale Structures)

Graph	Iter.	T	B	F	$B + F$	$(B + F)/T$
<i>General plane Graphs (Small to Medium)</i>						
graph_V5E6F3	14	4	11	3	14	3.50
graph_V5E7F3	6	2	5	1	6	3.00
graph_V5E7F4_1	6	2	5	1	6	3.00
graph_V5E7F4_2	1	1	1	0	1	1.00
graph_V5E7F4_3	6	2	5	1	6	3.00
graph_V6E7F3_1	46	20	27	19	46	2.30
graph_V6E7F3_2	58	24	35	23	58	2.42
graph_V6E7F3_3	46	20	27	19	46	2.30
graph_V6E7F3_4	46	20	27	19	46	2.30
graph_V7E9F3	14	9	6	8	14	1.56
graph_V7E11F5	12	7	6	6	12	1.71
graph_V7E14F9	2	2	1	1	2	1.00
graph_V8E10F4	94	40	55	39	94	2.35
graph_V8E15F8	10	5	6	4	10	2.00
<i>Large-Scale Instances</i>						
graph_V9E10F3	15,075	13,880	1,196	13,879	15,075	1.09
graph_V12E25F12	69	62	8	61	69	1.11
<i>Quadrilateral-based Graphs</i>						
quadrilateral_4quad	30	16	15	15	30	1.88
quadrilateral_7quad	162	72	91	71	162	2.25
quadrilateral_14quad	32,766	16,384	16,383	16,383	32,766	2.00

7 Conclusion and Future Work

We have presented the first algorithm for generating all triangulations of biconnected plane graphs with constant amortized time complexity. Our algorithm extends the elegant tree-of-triangulations framework of Parvez, Rahman, and Nakano [8] from triconnected to biconnected graphs by introducing safe root selection and local genealogical trees.

7.1 Key Achievements

1. **Theoretical Foundation:** We have provided rigorous proofs establishing:
 - Existence of at least two safe root vertices in every face (Theorem 1)
 - Completeness of the algorithm without duplications (Theorem 5)
 - Amortized $O(1)$ time complexity per triangulation (Theorem 8)
 - Linear $O(n)$ space complexity (Theorem 10)
2. **Practical Algorithm:** Efficient implementation suitable for real-world applications with extensive experimental validation
3. **Novel Techniques:** Safe root selection and global chord management strategies that can extend to other enumeration problems involving plane graph decompositions
4. **Broader Applicability:** Extension from triconnected to the more general class of biconnected plane graphs, significantly expanding the scope of graphs that can be processed efficiently

7.2 Theoretical Significance

This work makes several theoretical contributions:

- **Generalization:** We show that constant-time triangulation enumeration is achievable beyond the restrictive triconnected case
- **Multi-edge Resolution:** Our safe root concept provides a general strategy for handling chord conflicts in face-decomposed plane graphs
- **Amortized Analysis:** The detailed amortized analysis demonstrates that building time is dominated by output size even with complex face interactions

7.3 Future Directions

Promising extensions include: (1) extending to simply-connected plane graphs via block-cut tree decomposition, (2) parallel enumeration using lock-free data structures, (3) constrained triangulations for mesh generation applications, and (4) developing polynomial-time counting algorithms without explicit enumeration. Open problems include finding closed-form formulas for triangulation counts and uniform random sampling methods.

Acknowledgments

The author gratefully acknowledges the foundational work of Mohammad Tanvir Parvez, Md. Saidur Rahman, and Shin-ichi Nakano, whose elegant algorithm for triconnected plane graphs [8] inspired this research.

References

- [1] David Avis and Komei Fukuda. Reverse search for enumeration. *Discrete Applied Mathematics*, 65(1-3):21–46, 1996.
- [2] Mark de Berg, Otfried Cheong, Marc van Kreveld, and Mark Overmars. *Computational Geometry: Algorithms and Applications*. Springer, 3rd edition, 2008.
- [3] Reinhard Diestel. *Graph Theory*, volume 173 of *Graduate Texts in Mathematics*. Springer, 5th edition, 2017. See Theorem 2.3.1 (Maximal Outerplanar Graphs) for degree-2 vertex existence.
- [4] Ferran Hurtado and Marc Noy. Graph of triangulations of a convex polygon and tree of triangulations. *Computational Geometry*, 13(3):179–188, 1999.
- [5] Krzysztof Kozminski and Edwin Kinnen. Rectangular dualization and rectangular dissections. *IEEE Transactions on Circuits and Systems*, 32(8):771–781, 1985.
- [6] Zheng Li and Shin-ichi Nakano. Generating all triangulations of plane graphs. *Theoretical Computer Science*, 270(1-2):771–784, 2002.
- [7] Takao Nishizeki and Md Saidur Rahman. *Planar Graph Drawing*. World Scientific, 2004.
- [8] Mohammad Tanvir Parvez, Md. Saidur Rahman, and Shin-ichi Nakano. Generating all triangulations of plane graphs. *Journal of Graph Algorithms and Applications*, 15(3):457–482, 2011. See Section 4 for the generating chord concept used in local genealogical trees.

- [9] Ronald C Read. Every one a winner or how to avoid isomorphism search when cataloguing combinatorial configurations. *Annals of Discrete Mathematics*, 2:107–120, 1978.